

CS 3305A

# Multiprogramming

Lecture 7

Sept 30<sup>th</sup> 2019

# Multiprogramming

- ❑ Assume we have two processes  $p$  and  $q$
- ❑ Process  $p$  has an instruction that requires a read/write from/to disk
- ❑ Reading from disk is slow
- ❑ Why not have instructions from  $q$  execute while  $p$  is waiting?

# Multiprogramming

- ❑ **Multiprogramming** allows for the execution of multiple processes
- ❑ But only one process **active** at any time

# Why Multiprogramming?

- ❑ Operating systems allow for **interleaved execution**
  - ❑ On a single-processor system, no more than one process ever runs.
  - ❑ However, one process's instructions may be executed before the completion of the instructions from another process
- ❑ The **objective** is to have some process running at all times in order to maximize CPU utilization.

# Process Switching

- ❑ Current process executes an I/O operation
- ❑ OS needs to be able to suspend current process so that another process can execute
- ❑ This is referred to as **context switching**

# Process Switching

- ❑ OS needs to be able to suspend current process
- ❑ OS captures information about a process
- ❑ Information captured must be sufficient to restore the hardware to the same configuration it was in when the process was switched out.

# Characterizing a Process

- Each process is represented in the OS by a **process control block (PCB)** which contains all the state for a program in execution including (but not limited to):
  - Pointer to text, data etc. information
  - The program counter (PC) indicating the next instruction
  - Current values of the set of general-purpose registers
  - A set of operating system resources e.g., open files, network connections
  - Process identifier (PID)
  - Process priority (for scheduling purposes)
  - etc.

# Process Execution States

- ❑ As a process executes, it changes execution state
- ❑ The execution state of a process is defined in part by the current activity of the process
- ❑ A process may be in one of the following execution states:
  - ❑ **New**: The process is being created
  - ❑ **Ready**: The process is waiting to be assigned to a processor
  - ❑ **Running**: Instructions are being executed
  - ❑ **Waiting**: The process is waiting for some event to occur (such as an I/O completion or reception of signal)
  - ❑ **Exit**: The process has finished executing
- ❑ Only **one** process can be **running** on any processor at any instant
- ❑ **Many** processes may be **ready** and **waiting**



# Scheduling

- ❑ The purpose of multiprogramming is to have a process running at all times
- ❑ The objective of time sharing is to switch the CPU among processes so frequently that users can interact with each process
- ❑ The **process scheduler** selects an available process
- ❑ There may be multiple processes to select from

# Scheduling Queues

- ❑ As processes enter the system, they are put into a **job queue**, which consists of all processes in the system
- ❑ The processes that are residing in main memory and are ready and waiting to execute are kept on a list called the **ready queue**
- ❑ Queues are implemented using linked list

# Process Execution States

