

# CS 3305A


## System Calls

Lecture 6  
Sept 25<sup>th</sup> 2019

# Interface to the OS

- ❑ All operating systems have an interface to OS that is accessible by users/user programs
- ❑ We had a discussion of shells which allows a user to interface with the operating system through the command line
  - ❑ A second strategy is through a graphical user interface (GUI)
- ❑ We had seen how **system functions** (such as `fork()`) can communicate with OS

# A View of Operating System Services



# Interface to the OS

- ❑ **System calls** provide an interface to OS services
  - ❑ Program passes relevant information to OS
  - ❑ OS performs the service if
    - ❑ The OS is able to do so
    - ❑ The service is permitted for this program at this time
- ❑ System calls are typically written in C and C++
  - ❑ Tasks that require hardware to be accessed directly may be written using assembly language

# Application Programmer Interface (API)

- ❑ Programmers call a function (**system function**) in a library which invokes system calls
- ❑ The programmer only needs to understand the system function by understanding its parameters and results
- ❑ Example
  - ❑ Programmer API: `count = read(fd, buf, nbytes)`
  - ❑ System calls Used: `sys_read()`
  - ❑ System call code is part of the kernel (Core OS)

# Examples: Other System Calls

- ❑ Linux Examples:
  - ❑ `sys_fork`, `sys_pipe()`
- ❑ Note: We have been using system call loosely
  - ❑ Could be referring to the system function
  - ❑ System function and System call are two different entities
    - ❑ System function: used by user / programmer (API)
    - ❑ System call: Part of OS Kernel

# Some System Functions For Process Management

## Process management

Call	Description
<code>pid = fork( )</code>	Create a child process identical to the parent
<code>pid = waitpid(pid, &amp;statloc, options)</code>	Wait for a child to terminate
<code>s = execve(name, argv, environp)</code>	Replace a process' core image
<code>exit(status)</code>	Terminate process execution and return status

# Some System Functions For File Management

## File management

Call	Description
<code>fd = open(file, how, ...)</code>	Open a file for reading, writing or both
<code>s = close(fd)</code>	Close an open file
<code>n = read(fd, buffer, nbytes)</code>	Read data from a file into a buffer
<code>n = write(fd, buffer, nbytes)</code>	Write data from a buffer into a file
<code>position = lseek(fd, offset, whence)</code>	Move the file pointer
<code>s = stat(name, &amp;buf)</code>	Get a file's status information



# Some System Functions For Directory Management

## Directory and file system management

Call	Description
s = mkdir(name, mode)	Create a new directory
s = rmdir(name)	Remove an empty directory
s = link(name1, name2)	Create a new entry, name2, pointing to name1
s = unlink(name)	Remove a directory entry
s = mount(special, name, flag)	Mount a file system
s = umount(special)	Unmount a file system

# APIs

- ❑ Let's say that a user program has the following line of code: `read(fd, buf, nbytes)`
- ❑ This program needs the operating system to access the file and read from it.
- ❑ Some issues to be addressed:
  - ❑ How are parameters passed?
  - ❑ How are results provided to the user program?
  - ❑ How is control given to the system call and the operating system?

# System Call Parameter Passing


- ❑ Three general methods used to pass parameters to the OS
  - ❑ **Registers:** Pass the parameters in registers
    - ❑ In some cases, there may be more parameters than registers
  - ❑ **Block:** Parameters stored in a **block**, or table, in memory, and address of block passed as a parameter in a register
    - ❑ This approach taken by Linux and Solaris
  - ❑ **Stack:** Parameters placed, or pushed, onto the stack by the program and popped off the stack by the operating system
- ❑ Block and stack methods do not limit the number or length of parameters being passed

# Linux: Parameter passing

- ❑ System calls with fewer than 6 parameters passed in registers
- ❑ If 6 or more arguments
  - ❑ Pass pointer to block structure

# System Call Table

- ❑ A system call number is associated with each system call
- ❑ The OS maintains a **system call handler table** which is indexed according to the system call numbers
- ❑ Entry in table points to code



# System Calls and Traps

- ❑ **TRAP** switches CPU to supervisor (kernel) mode
  - ❑ The state of the user process is saved so that the OS instructions needed can be executed (**system call**)
  - ❑ When the system handler finishes execution then the user process can execute

# Making a System Call

- ❑ System function call:

`count = read(fd,buffer,length)`

- ❑ Step 1: The input parameters are passed into registers or to a block
- ❑ Step 2: TRAP (execution of system call) is executed
  - ❑ The state of the user process is saved
  - ❑ System call number for read() is sent to system call handler
  - ❑ This code/number tells the OS what system call handler (kernel code) to execute
  - ❑ This causes a switch from the user mode to the kernel mode

# Making a System Call

- ❑ Step 3: System call handler code is executed
- ❑ Step 4: After execution control returns to the library procedure (system function)



# System Call

- ❑ The system call handler will have to actually wait for data from the disk
- ❑ Reading data from disk is much slower than memory
- ❑ We do not want the CPU to be idle while waiting for data from the disk
  - ❑ Most operating systems allow for another executing program to use the CPU
  - ❑ This is called **multiprogramming** - more later
- ❑ How does a process find out about reading being completed?
  - ❑ Interrupt